

HW2

一、环境配置

1.1 conda环境创建

```
conda create -n gym2 python=3.8
```

1.2 安装依赖

numpy, matplotlib, tqdm按照文档安装即可

gym根据文档安装0.24.0版本，但是运行代码的时候报错reset返回值是一个int而不是tuple，推测是API版本过低，gym升级到最新版即可

二、Blank fill

2.1 update_q_value

```
def update_q_value(q, q_next, reward, alpha, gamma):  
    """  
    TODO:  
    Please fill in the blank for variable 'td_target'  
    according to the definition of the TD method  
    """  
    td_target = reward + gamma * q_next  
    return q + alpha * (td_target - q)
```

sarsa和Q-learning都是表格型时序差分方法，策略评估为更新状态-动作值函数：

SARSA：

$$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha(R_{t+1} + \gamma Q(s_{t+1}, a_{t+1}) - Q(s_t, a_t))$$

Q-learning：

$$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha(R_t + \gamma \max_{a'} Q(s_{t+1}, a'_{t+1}) - Q(s_t, a_t))$$

其中reward为奖励，gamma为折扣因子，填入即可

2.2 sarsa

```
"""  
    TODO:  
    Please fill in the blank for variable 'next_q_value'  
    according to the definition of the SARSA algorithm  
    """  
    next_q_value = q_values[next_state, next_action]  
    q_values[state, action] = update_q_value(  
        q_values[state, action],  
        0 if done else next_q_value,  
        reward, config.alpha, config.gamma)
```

根据SARSA的on-policy时序差分算法:

Sarsa: An on-policy TD control algorithm

```
Initialize  $Q(s, a), \forall s \in \mathcal{S}, a \in \mathcal{A}(s)$ , arbitrarily, and  $Q(\text{terminal-state}, \cdot) = 0$ 
Repeat (for each episode):
  Initialize  $S$ 
  Choose  $A$  from  $S$  using policy derived from  $Q$  (e.g.,  $\epsilon$ -greedy)
  Repeat (for each step of episode):
    Take action  $A$ , observe  $R, S'$ 
    Choose  $A'$  from  $S'$  using policy derived from  $Q$  (e.g.,  $\epsilon$ -greedy)
     $Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma Q(S', A') - Q(S, A)]$ 
     $S \leftarrow S'; A \leftarrow A';$ 
  until  $S$  is terminal
```

此处选择下一个状态的Q值填入即可

2.3 q_learning

```
"""
    TODO:
    Please fill in the blank for variable 'next_q_value'
    according to the definition of the Q-learning algorithm
    """

next_q_value = np.max(q_values[next_state])
q_values[state, action] = update_q_value(
    q_values[state, action],
    0 if done else next_q_value,
    reward, config.alpha, config.gamma)
```


对于Q-learning算法, next_q_value是下一状态所有动作的最大Q值。

三、Test


3.1 CliffWalking

测试结果如下:

```
(gym2) PS E:\ZJU2024-MARL-JUEWU> python .\ex2_cliff-walking.py
Discrete(4)
Discrete(48)
100%|████████████████████████████████████████████████████████████████████████████████| 10000/10000 [00:05<00:00, 1729.27it/s]
Best Average Reward over 100 Episodes: -13
100%|████████████████████████████████████████████████████████████████████████████████| 10000/10000 [00:06<00:00, 1585.93it/s]
Best Average Reward over 100 Episodes: -17
```



可以看到结果和教材给出的大体相同：




3.2 FrozenLake

修改代码如下：

```
env = gym.make('FrozenLake-v1')
config = Configs(env, max_timestep=200, num_episode=40000, plot_every=100)
```

测试结果如下：



```
(gym2) PS E:\ZJU2024-MARL-JUEWU> python .\ex2_cliff-walking.py
Discrete(4)
Discrete(16)
100% | 40000/40000 [00:33<00:00, 1185.53it/s]
Best Average Reward over 100 Episodes: 1.0
100% | 40000/40000 [00:25<00:00, 1583.88it/s]
Best Average Reward over 100 Episodes: 1.0
```