

# HW4

## 一、环境配置

### 1.1 依赖安装

沿用HW3的环境即可

`pip install gymnasium[Atari]` `pip install gymnasium[accept-rom-license]` 下载安装 gymnasium 以及 Atari 环境和 rom

## 二、Blank fill

### 2.1 QNetwork

```
class QNetwork(nn.Module):
    def __init__(self, env):
        super().__init__()
        '''
        定义一个三层神经网络，其中
        输入层：全连接层： 状态空间作为输入，256作为输出，ReLU激活函数
        中间层：全连接层：256输入256输出，ReLU激活
        输出层：全连接层：256输入，动作空间输出
        '''
        self.fc1 = nn.Linear(env.observation_space.shape[0], 256)
        self.fc2 = nn.Linear(256, 256)
        self.fc3 = nn.Linear(256, env.action_space.n)

    def forward(self, x):
        '''
        前向计算
        '''
        x = F.relu(self.fc1(x))
        x = F.relu(self.fc2(x))
        x = self.fc3(x)
```

按照提示填入即可

### 2.2 define Network

```
'''
    定义两个网络，分别为q_network, 以及target_network;
    定义优化器，训练网络
    '''
    q_network = QNetwork(env).to(device)
    target_network = QNetwork(env).to(device)
    target_network.load_state_dict(q_network.state_dict())
    target_network.eval()

    optimizer = optim.Adam(q_network.parameters(), lr=learning_rate)
```

初始化两个相同网络，分别作为评估网络和目标网络

## 2.3 epsilon-greedy

```
'''
    实现epsilon-greedy算法, epsilon为给定超参
    '''
    if random.random() < epsilon:
        actions = env.action_space.sample()
    else:
        with torch.no_grad():
            obs_tensor = torch.tensor(obs, device=device,
dtype=torch.float32).unsqueeze(0)
            actions = q_network(obs_tensor).argmax(dim=1).item()
```

$\epsilon$ 贪心 ( $\epsilon$ -greedy) 策略

$\epsilon$ -greedy  $\pi(s)$

$$= \begin{cases} \operatorname{argmax}_a q_\pi(s, a), & \text{以 } 1 - \epsilon \text{ 的概率} \\ \text{随机的 } a \in A, & \text{以 } \epsilon \text{ 的概率} \end{cases}$$

## 2.4 td\_target

```
'''
    计算td_target
    Q(s,a) (old_val)
    '''
    with torch.no_grad():
        target_max = target_network(next_buffer_obs).max(dim=1,
keepdim=True)[0]
        td_target = rew + cont * gamma * target_max
        old_val = q_network(buffer_obs).gather(1, act)
```

使用目标网络计算TD:

$$y = r + \gamma Q_{\omega'}(s', \operatorname{argmax}_{a'} Q_{\omega'}(s', a'))$$

## 2.5 运行结果 (时间步 $10^6$ ) :

td_loss: 0.000815562903881073	q_values: 1.4344158172607422	step: 998100, avg_rewards: 22.93
td_loss: 0.0005887472070753574	q_values: 1.4041332006454468	step: 998200, avg_rewards: 22.93
td_loss: 0.0013498843181878328	q_values: 1.377229928970337	step: 998300, avg_rewards: 22.93
td_loss: 0.00135188945569098	q_values: 1.459467887878418	step: 998400, avg_rewards: 22.93
td_loss: 0.0018515526317059994	q_values: 1.4182753562927246	step: 998500, avg_rewards: 22.93
td_loss: 0.0007640982512384653	q_values: 1.3815784454345703	step: 998600, avg_rewards: 22.93
td_loss: 0.0014886329881846905	q_values: 1.4088695049285889	step: 998700, avg_rewards: 22.93
td_loss: 0.0007334256661124527	q_values: 1.3678758144378662	step: 998800, avg_rewards: 22.93
td_loss: 0.0019527716794982553	q_values: 1.4145925045013428	step: 998900, avg_rewards: 22.93
td_loss: 0.0006778707029297948	q_values: 1.4363003969192505	step: 999000, avg_rewards: 22.93
td_loss: 0.0013655954971909523	q_values: 1.402050495147705	step: 999100, avg_rewards: 22.93
td_loss: 0.0009839077247306705	q_values: 1.4100652933120728	step: 999200, avg_rewards: 22.93
td_loss: 0.0010703383013606071	q_values: 1.3783280849456787	step: 999300, avg_rewards: 22.93
td_loss: 0.00041139504173770547	q_values: 1.3635199069976807	step: 999400, avg_rewards: 22.93
td_loss: 0.0010131148155778646	q_values: 1.4151065349578857	step: 999500, avg_rewards: 22.94
td_loss: 0.0013535680482164025	q_values: 1.408816933631897	step: 999600, avg_rewards: 22.94
td_loss: 0.0004961451049894094	q_values: 1.426405906677246	step: 999700, avg_rewards: 22.94
td_loss: 0.0011145444586873055	q_values: 1.4043384790420532	step: 999800, avg_rewards: 22.94
td_loss: 0.0012148329988121986	q_values: 1.384615182876587	step: 999900, avg_rewards: 22.94

## 三、DDQN

### 3.1 td\_target

```
with torch.no_grad():
    next_actions = q_network(next_buffer_obs).argmax(dim=1,
keepdim=True)
    target_q_values = target_network(next_buffer_obs).gather(1,
next_actions)
    td_target = rew + cont * gamma * target_q_values
    old_val = q_network(buffer_obs).gather(1, act)[:-1]
```

和DQN不同的是，action和目标价值的估计是通过两个网络来进行的

### 3.2 运行结果

td_loss: 0.0013962259981781244	q_values: 1.0258598327636719	step: 997700, avg_rewards: 22.46
td_loss: 0.0018020548159256577	q_values: 1.0453356504440308	step: 997800, avg_rewards: 22.46
td_loss: 0.001169784227386117	q_values: 1.055753469467163	step: 997900, avg_rewards: 22.46
td_loss: 0.0009516063728369772	q_values: 1.0622200965881348	step: 998000, avg_rewards: 22.46
td_loss: 0.0018792767077684402	q_values: 1.134387493133545	step: 998100, avg_rewards: 22.46
td_loss: 0.001369684236124158	q_values: 1.0754268169403076	step: 998200, avg_rewards: 22.46
td_loss: 0.0009112475090660155	q_values: 1.070970058441162	step: 998300, avg_rewards: 22.46
td_loss: 0.0010699429549276829	q_values: 1.0656156539916992	step: 998400, avg_rewards: 22.46
td_loss: 0.0011589701753109694	q_values: 1.033522367477417	step: 998500, avg_rewards: 22.46
td_loss: 0.0009377492242492735	q_values: 1.040408730506897	step: 998600, avg_rewards: 22.46
td_loss: 0.002161842305213213	q_values: 1.0578950643539429	step: 998700, avg_rewards: 22.46
td_loss: 0.0015440761344507337	q_values: 1.0505661964416504	step: 998800, avg_rewards: 22.46
td_loss: 0.0012774900533258915	q_values: 1.0797215700149536	step: 998900, avg_rewards: 22.46
td_loss: 0.0014378158375620842	q_values: 1.0943503379821777	step: 999000, avg_rewards: 22.46
td_loss: 0.001382244867272675	q_values: 1.054526925086975	step: 999100, avg_rewards: 22.46
td_loss: 0.0013589609880000353	q_values: 1.0118794441223145	step: 999200, avg_rewards: 22.46
td_loss: 0.0008017393993213773	q_values: 1.0452497005462646	step: 999300, avg_rewards: 22.46
td_loss: 0.002535087987780571	q_values: 1.001410961151123	step: 999400, avg_rewards: 22.46
td_loss: 0.0009350934997200966	q_values: 1.044939398765564	step: 999500, avg_rewards: 22.46
td_loss: 0.0024224319495260715	q_values: 1.0549191236495972	step: 999600, avg_rewards: 22.46
td_loss: 0.0012535869609564543	q_values: 1.0680443048477173	step: 999700, avg_rewards: 22.46
td_loss: 0.0012248046696186066	q_values: 1.0616261959075928	step: 999800, avg_rewards: 22.46
td_loss: 0.0012793298810720444	q_values: 1.0420677661895752	step: 999900, avg_rewards: 22.46

可以看到，DDQN减少了过高估计的偏差